



TUBAF

Die Ressourcenuniversität.
Seit 1765.

FAKULTÄT FÜR MATHEMATIK UND INFORMATIK

TECHNISCHE UNIVERSITÄT BERGAKADEMIE FREIBERG

HPC Project

Disease Spread Simulation on a 2D Grid using open MPI and openMP

Author: Omkar Kolekar
Submission Date: 26 June 2024



Contents

1	Introduction	1
1.1	Chapters	1
2	Problem Description	2
2.1	Grid States	2
2.2	Rules for State Changes	2
3	Algorithm and User Manual	3
3.1	User Manual	3
4	Validation and testing	4
4.1	Visualization of output	4
4.2	Testing	5
5	Strong and Weak Scability Test	6
5.1	Strong scaling test	6
5.2	Inference	6
5.2.1	Without multi threading	6
5.2.2	With multi threading	6
5.3	Weak scaling test	8
5.3.1	Inference	10
6	Conclusion	11
	List of Figures	12
	List of Tables	13

1 Introduction

Infectious diseases persist as a formidable challenge to global health, underscoring the urgent need for advanced computational methods to study their transmission dynamics comprehensively. Addressing this need, the current simulation harnesses the robust capabilities of parallel programming frameworks, specifically OpenMPI and OpenMP, implemented within the C++ programming language. These frameworks facilitate the modeling of infectious disease spread within a simulated population by enabling efficient multi-threading and distributed computing. Through the utilization of OpenMPI and OpenMP, the simulation aims to investigate the effectiveness and scalability of parallel computing when applied to large-scale epidemic scenarios. Furthermore, a Python script is integrated into the workflow to provide visualization capabilities, generating intuitive graphical representations of the simulation results. This study is dedicated to evaluating the computational advantages and practical applications of parallel programming in the context of epidemiological simulations.

Parallel computing revolutionizes the speed and efficiency of computational tasks by concurrently utilizing multiple processors or computing resources. Unlike traditional serial computing, where tasks are processed sequentially, parallel computing allows for simultaneous execution, leading to significant performance enhancements and scalability. This approach is essential for tackling intensive computations in fields such as scientific simulations, big data analytics, AI development, and real-time processing, where speed, efficiency, and scalability are paramount for advancing technological capabilities and driving innovation.

1.1 Chapters

The report is organized into six chapters as follows.

- Chapter 1:** Introduction to the report.
- Chapter 2:** Problem statement description.
- Chapter 3:** Algorithm description and User manual.
- Chapter 4:** Validation and testing.
- Chapter 5:** Strong and weak scalability test.
- Chapter 6:** Conclusion.

2 Problem Description

In this project, a parallel simulation of disease spread on a grid is to be implement, inspired by SIR (Susceptible, Infected, Recovered) or SEIR (Susceptible, Exposed, Infected, Recovered) models. This simulation will use MPI for spatial partitioning of the grid and OpenMP for parallel processing within each grid segment.

2.1 Grid States

Each cell in the grid can have one of the following states:

- **Susceptible (Healthy):** Can be infected.
- **Infected:** Can infect other cells and will become healthy after a certain time.
- **Recovered (Immune):** Temporarily immune after recovering from infection.

2.2 Rules for State Changes

1. **Infection:** A healthy cell adjacent to an infected cell becomes infected with a probability p (infection probability) in the next time step.
2. **Recovery and Immunity:** An infected cell has a probability q (recovery probability) to become healthy in the next time step. After recovery, it becomes immune for t time steps. After the immunity period, the cell becomes susceptible again.

3 Algorithm and User Manual

The below flow chart gives an overview of the working algorithm.

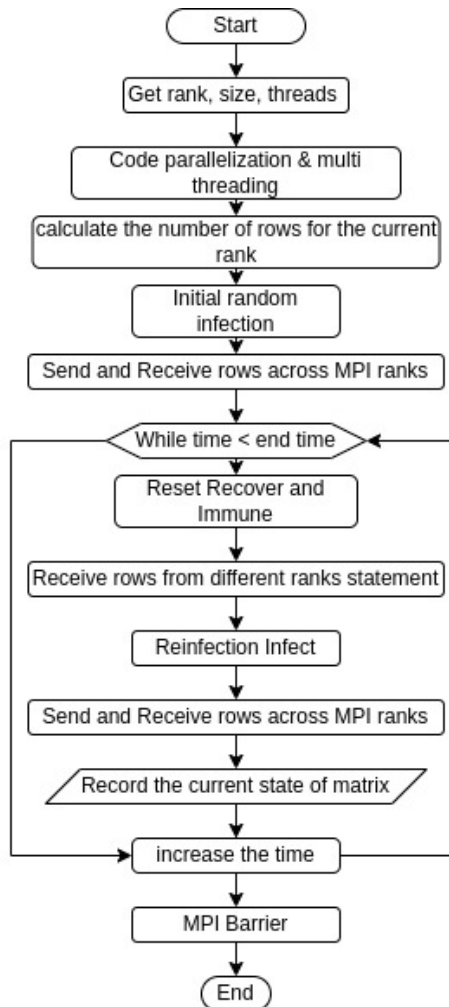


Figure 3.1: Algorithm flow chart


3.1 User Manual

The user should ensure that the Linux system has openmpi and openmp installed.


4 Validation and testing

4.1 Visualization of output

The below few pictures are snaps from first two time steps. The size of the matrix was 75x75.




(a) Infected




(b) Recovered

Figure 4.1: Infection state at time $t = 0$



(a) Infected



(b) Recovered

Figure 4.2: Infection state at time $t = 1$

4.2 Testing

Since I have chosen row major distribution of matrix across all the MPI ranks, the testing in MPI is based on the boundary rows of the matrices.

The software is programmed in such a way that each rank will write the state of the matrix to a text file `matrix_output_rank_X.txt` (where: 'X' rank number) irrespective of the number of ranks. For a new time step a new matrix will be appended under the previous matrix in the same text file.

Using the `cat` command in the linux bash terminal, following pictures were obtained. I have marked the important rows of the respective rank. The red rectangle represent the state of the row at time $t = 0$. The green rectangle represent the state of the row at time $t = 1$. The boundary rows show no irregularities in all the ranks, which ensures the proper implementation of the MPI parallelization.

```
The Simulation started with total ranks = 3 and total threads = 2
Printing the initial stage of the matrix
Printing the stage of the matrix at time t = 1
The time taken is 0.00871615 for 3 number of ranks
The time taken is 0.00876644 for 3 number of ranks
The Simulation completed with no errors.
All the necessary files have been saved successfully.
The time taken is 0.00891394 for 3 number of ranks
[ok28rune@mfatnode006 Final_Project]$ ls
EpidemiSimPar.cpp  Visualize.py  compilefile.sh  matrix_output_rank_0.txt  matrix_output_rank_2.txt  onlyrun.sh
FinalProject.cpp  a.out          load_the_mpi_file.sh  matrix_output_rank_1.txt  noderunning.sh          run_comp.sh
[ok28rune@mfatnode006 Final_Project]$ cat matrix_output_rank_0.txt
1; 0.1; 0; 0; 0.1; 1; 0.1; 0; 0; 0; 0; 0; 0; 0; 0;
0.1; 0.1; 0; 0; 0.1; 0.1; 0.1; 0; 0; 0; 0; 0; 0; 0;
0; 0; 0; 0; 0; 0; 0.1; 0.1; 0.1; 0; 0; 0; 0; 0;
0; 0; 0; 0; 0; 0.1; 0.1; 0.2; 1; 0.1; 0; 0.1; 0.1; 0;
0; 0; 0.1; 0.1; 0.1; 0.1; 1; 0.2; 0.1; 0.1; 0.1; 0.2; 1; 0.1; 0;
1; 0.1; 0; 0; 0; 0; 2; 0; 0; 0; 0; 0; 0; 0; 0;
0.1; 0.1; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0;
0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0;
0; 0; 0; 0; 0; 0.1; 0.1; 0.1; 2; 0; 0.1; 0.2; 0.2; 0.1; 0;
0; 0; 0; 0; 0; 0.1; 1; 0.1; 0; 0; 0.1; 1; 1; 0.1; 0;
[ok28rune@mfatnode006 Final_Project]$ cat matrix_output_rank_1.txt
0; 0; 0.1; 1; 0.1; 0.2; 0.2; 0.2; 0; 0; 0.1; 1; 0.2; 0.1; 0;
0; 0; 0.1; 0.2; 0.2; 0.2; 1; 0.1; 0; 0; 0.1; 0.1; 0.1; 0; 0;
0; 0; 0; 0.1; 1; 0.2; 0.1; 0.1; 0; 0; 0; 0; 0; 0; 0;
0; 0; 0; 0.1; 0.1; 0.1; 0; 0.1; 0.1; 0.1; 0; 0; 0; 0; 0;
0; 0; 0; 0; 0; 0; 0.1; 1; 0.1; 0; 0.1; 0.1; 0.1; 0;
0; 0; 0; 2; 0; 0.1; 0.1; 0.1; 0; 0; 0.1; 2; 0.2; 0.1; 0;
0; 0; 0; 0; 0; 0; 2; 0; 0; 0; 0; 0; 0; 0; 0;
0; 0; 0; 0; 2; 0; 0.1; 0.2; 0.3; 0.3; 0.2; 0.1; 0; 0; 0;
0; 0; 0; 0; 0; 0.2; 1; 1; 1; 1; 0.1; 0; 0; 0; 0;
0; 0; 0; 0; 0; 0.3; 1; 1; 0.4; 0.2; 0.1; 0; 0; 0; 0;
[ok28rune@mfatnode006 Final_Project]$ cat matrix_output_rank_2.txt
0; 0; 0; 0; 0; 0; 0.1; 0.2; 0.2; 0.1; 0; 0.1; 1; 0.1; 0;
0.1; 0.1; 0.1; 0; 0; 0; 0.1; 1; 0.2; 0.1; 0.1; 0.2; 0.2; 0.1; 0;
0.1; 1; 0.1; 0; 0; 0; 0.1; 0.2; 1; 0.1; 0.1; 1; 0.1; 0; 0;
0.1; 0.1; 0.1; 0; 0; 0; 0.1; 0.1; 0.1; 0.1; 0.1; 0; 0; 0;
0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0;
0; 0.1; 0.1; 0.1; 0; 0; 0.3; 1; 0.4; 0.1; 0; 0; 2; 0; 0;
0.2; 0.3; 1; 0.1; 0; 0; 0.2; 1; 0.2; 0; 0; 0; 0; 0; 0;
1; 1; 0.3; 0.2; 0; 0; 0.1; 0.1; 2; 0; 0; 2; 0; 0; 0;
0.2; 0.3; 1; 0.1; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0;
0; 0.1; 0.1; 0.1; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0; 0;
```

Figure 4.3: Boundary row testing result

5 Strong and Weak Scalability Test

Scaling tests are employed to assess the parallel performance of a computational code, with the objective of optimizing the performance of an application or system. These tests enable the determination of the most efficient allocation of computational resources for a given problem. It is noteworthy that during scaling tests, the grid is not printed to the output file. Adjustments to the grid size and the number of iterations are implemented through code modifications, rather than being provided as user inputs for each trial.

5.1 Strong scaling test

A strong scaling test is a test in which the problem size remains constant while the computing power is variable. The size of the grid was 1000 X 1000 and the number of time steps were 50. The number of cores (Processors) were varied from 1 to 40 and the threads per processor was also varied from 1 to 2 to observe the effect of multi-threading. The table 5.1 compares the the effect of multi threading per core on reduction in computational time as the number of cores (processors) increases. The figure 5.2 is a visual representation of table 5.1.

5.2 Inference

5.2.1 Without multi threading

The figure 5.1a illustrates the reduction in computational time as the number of cores (processors) increases, each utilizing a single thread per core.

- With increase in the number of cores (processors) the computational time does decrease up to a certain limit.
- Beyond this limit there is no overall drop in the computational time as the overhead increases

5.2.2 With multi threading

The figure 5.1b illustrates the reduction in computational time as the number of cores (processors) increases, each utilizing a two threads per core.

- There is no significant reduction in the computational time observed due to multi threading.

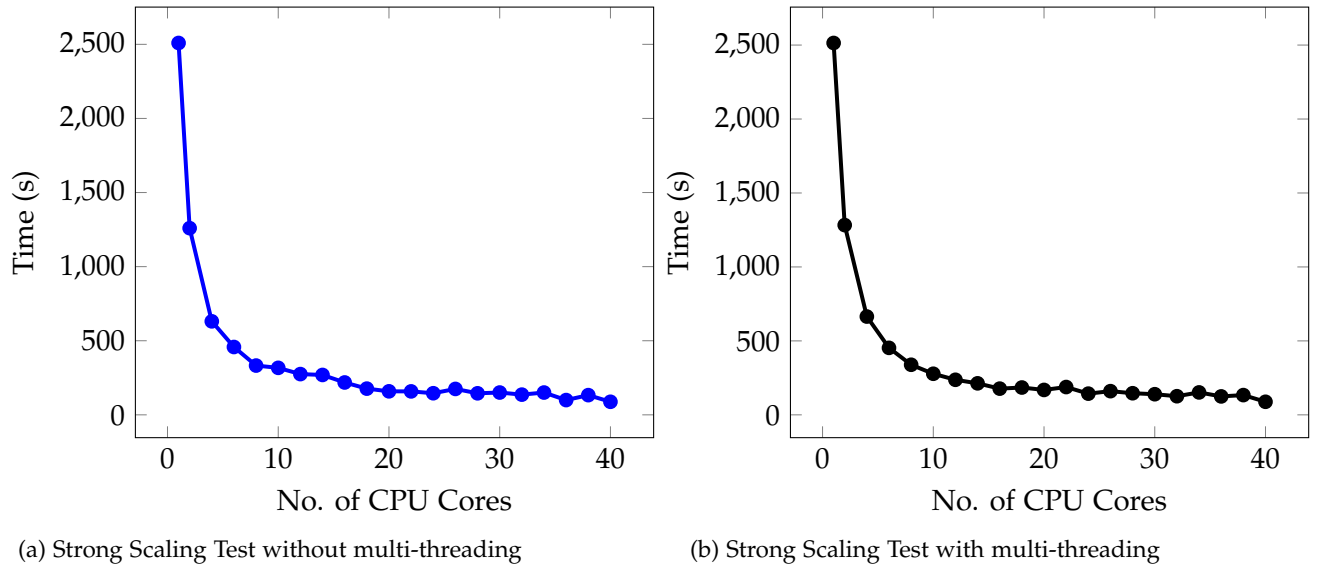


Figure 5.1: Comparison of strong scaling tests

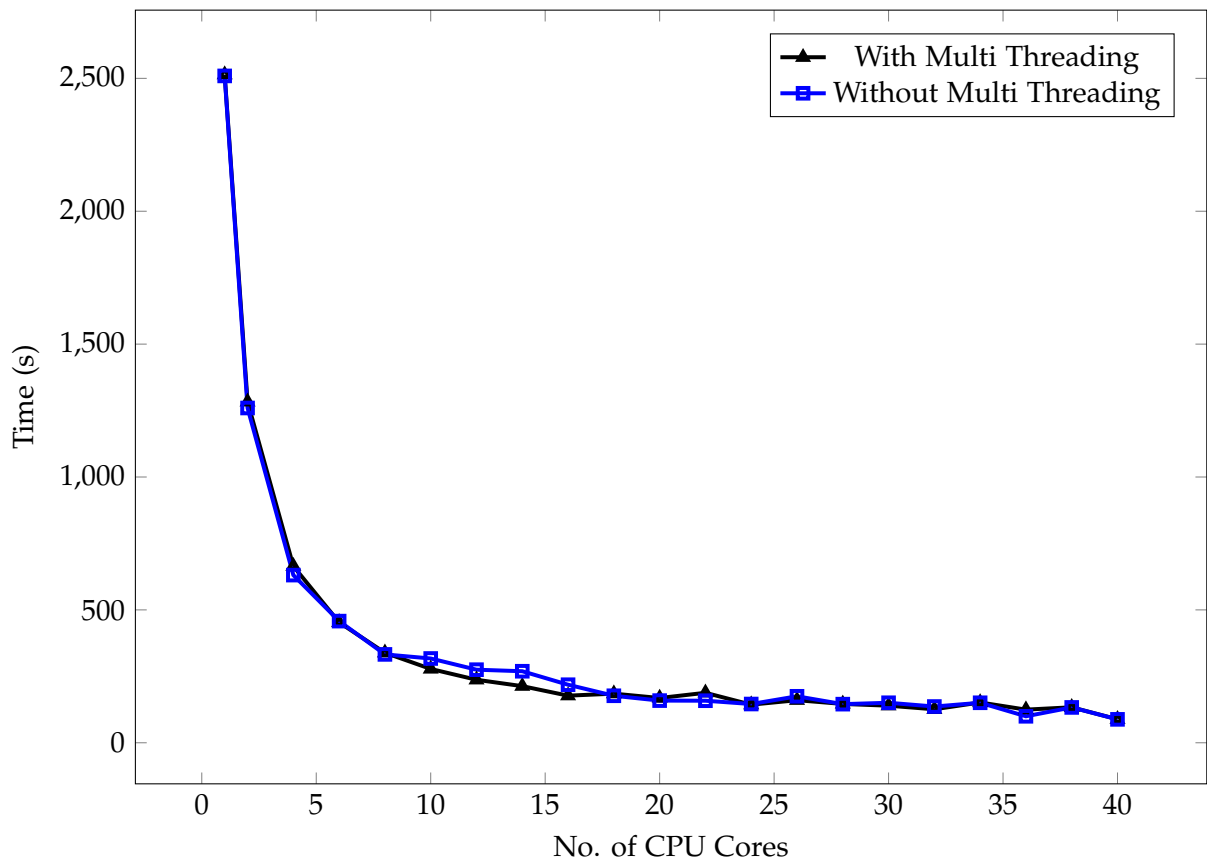


Figure 5.2: Effect of multi-threading on Strong Scaling Test

No. of Cores	Time (s) W/o Multi Threading	Time (s) Multi Threading
1	2509.17	2514.08
2	1259.65	1282.93
4	630.918	664.55
6	457.772	452.93
8	332.514	338.638
10	317.133	277.815
12	274.985	237.062
14	269.151	213.066
16	218.431	177.263
18	176.879	184.809
20	158.694	168.437
22	158.261	188.026
24	145.816	143.395
26	174.889	160.205
28	145.322	146.166
30	150.56	139.528
32	136.389	126
34	150.501	151.664
36	99.67	124.709
38	132.98	133.567
40	88.1205	88.3432

Table 5.1: Strong Scaling Test.

5.3 Weak scaling test

Weak scaling tests assess how well a program can manage larger problems by increasing resources while maintaining the workload per processing unit constant. This capability is crucial as it demonstrates the program's ability to handle larger problems effectively when sufficient computational resources are available. In this test the number of cores were doubled from 1 to 32 for each run. The test is conducted for three different load cases i.e. $10E4$ cells, $10E5$ cells and $5E5$ cells per core respectively. For all trials the time steps were kept constant. The graph 5.3 was obtained from the data in the table 5.2 which summarises the wall time required to complete the computation under different load cases.

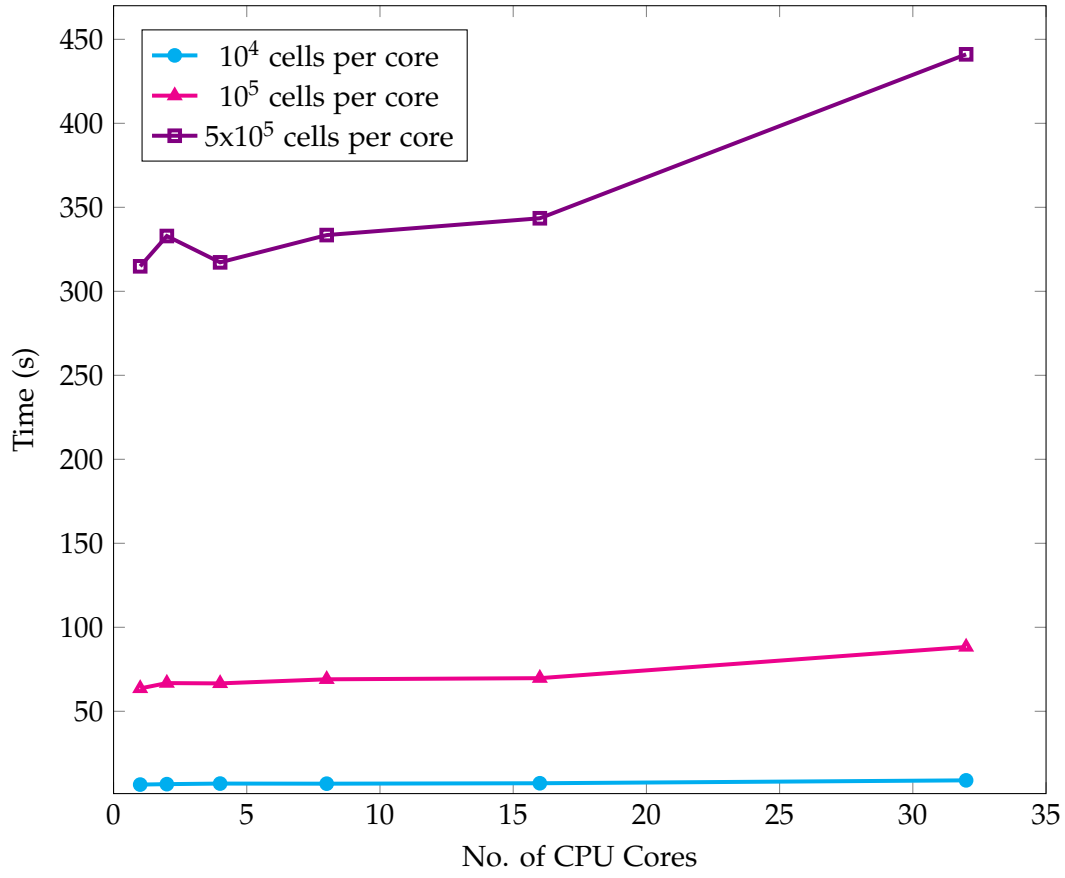


Figure 5.3: Weak Scaling Test

Number of cores	Average wall time (s)		
	10E4 cells	10E5 cells	5E5 cells
1	6.36437	63.6063	314.873
2	6.62011	66.7798	332.909
4	6.96612	66.5947	317.231
8	6.90212	69.0349	333.472
16	7.15825	69.7105	343.477
32	8.8799	88.2795	441.062

Table 5.2: Weak Scaling Test Data

5.3.1 Inference

- The curve with $10E4$ cells per core exhibits an almost flat slope, which is indicative of an ideal scenario for the weak scaling test. However, this is not the ideal load to measure the weak scaling test as the wall time is below 50 s.
- The curve with $10E5$ cells per core is almost flat as well which is an ideal behaviour.
- The slope of the last curve with $5E5$ cells per core deviate significantly from the ideal behaviour. The curve has a positive slope.

6 Conclusion

In this project, the spread of disease within a community was simulated. The simulation was programmed in C++ utilizing MPI and OpenMP. The program outputs a matrix in the form of text files, with the number of text files generated corresponding to the number of threads used. Race conditions within the threads were avoided through the proper placement of barriers. Strong scaling tests yielded good results, demonstrating a significant reduction in computation time with parallelization. However, the effect of multi-threading was found to be insignificant, likely due to the increase in overhead. Weak scaling tests indicated that the program performs well for approximately 5,000,000 cells per core. In conclusion, the simulation effectively leverages parallel computing to model the spread of disease, significantly improving computational efficiency

List of Figures

3.1	Algorithm flow chart	3
4.1	Infection state at time $t = 0$	4
4.2	Infection state at time $t = 1$	4
4.3	Boundary row testing result	5
5.1	Comparison of strong scaling tests	7
5.2	Effect of multi-threading on Strong Scaling Test	7
5.3	Weak Scaling Test	9

List of Tables

5.1	Strong Scaling Test.	8
5.2	Weak Scaling Test Data	9